Kleinanzeigen Internal AI Policy

Document ID: KA-AI-POL-001 | **Version:** 1.0 | **Effective:** 1 October 2026 | **Owner:** AI Governance Board | **Review cycle:** every 12 months

Related documents: KA-AI-TOOLS-002 Approved Tools List; KA-BV-AI-2026 Works Council Agreement on AI and Digital Tools (Betriebsvereinbarung); KA-AI-PROJ-004 Past AI and Process Projects.

FICTIONAL DEMO DOCUMENT. Written for an Ironhack bootcamp project. It is not an actual Kleinanzeigen policy and contains no real company data.

1. Purpose and Scope

This policy sets the rules for using artificial intelligence (AI) at Kleinanzeigen. Its goal is to help employees use AI to remove friction from their daily work, quickly and safely, while protecting customers, employees and the company. It applies to all employees, contractors and temporary staff, in all departments and locations, and to every AI system used for Kleinanzeigen work: internal tools, approved external tools and any new tool that is being tested.

The policy follows the EU AI Act (Regulation (EU) 2024/1689), the GDPR (DSGVO), the German Works Constitution Act (Betriebsverfassungsgesetz, BetrVG) and the Works Council Agreement KA-BV-AI-2026. If this policy and the law differ, the law applies.

2. Roles and Responsibilities

To keep reviews fast, **one body owns every AI review: the AI Governance Board**. Employees do not need to find the right expert themselves. All edge cases go to one review channel, and the Board decides together.

Role	Responsibility
AI Governance Board (single review owner)	Reviews all Amber and Red cases, decides approve / approve with conditions / reject, keeps this policy and the Approved Tools List up to date. Members: Head of AI Enablement (CPO organisation, chair), Data Protection Officer (Datenschutzbeauftragte/r), Information Security Lead, a Works Council (Betriebsrat) representative and Legal Counsel.
Review channel	Slack channel **#ai-governance-review**. All escalations are posted here. The Board meets every Tuesday to review open cases together; urgent cases can be decided asynchronously by the chair plus the Data Protection Officer.
Operational Effectiveness and Enablement (OEE) team	Helps employees describe process problems, supports adoption of approved tools, collects feedback.
Tool owner	Named owner of each tool in KA-AI-TOOLS-002. Answers questions, manages access, reports incidents.
Employees	Use AI only as this policy allows, start with internal tools, never enter data above a tool's approved data class, report incidents.

3. Definitions

Term	Meaning in this policy
AI system	Software that generates content, predictions, recommendations or decisions from input data, including generative AI (chatbots, writing and coding assistants) and machine learning models.
Internal tool	An AI tool built or hosted by Kleinanzeigen, listed as "internal" in KA-AI-TOOLS-002.
Approved external tool	A third-party tool listed as "approved" in KA-AI-TOOLS-002 for a stated data class.
Pilot tool	A tool under test with a limited user group. Use only inside the pilot scope.
Not approved tool	Any tool that is not listed, or listed as "not approved". Treated as Red until reviewed.
Personal data	Any information about an identifiable person (customers, users, employees, candidates).
Employee data	Personal data about Kleinanzeigen employees, including performance, behaviour, attendance, location and communication data.
Human in the loop	A named person checks and approves AI output before it is used or sent.

4. Data Classification for AI Use

Every AI use starts with one question: **what is the highest data class I will put into the tool?** A tool may only receive data up to the class it is approved for in KA-AI-TOOLS-002.

Class	Examples	AI rule
D1 Public	Published help pages, public listings, press releases	Any approved tool, internal or external.
D2 Internal	Process descriptions, internal guides, aggregated KPIs, meeting notes without names	Internal tools and external tools approved for D2 or higher.
D3 Confidential	Financials, contracts, product roadmaps, security details, separation programme plans	Internal tools only, unless a tool is explicitly approved for D3.
D4 Personal data	Customer messages, user IDs, seller contact details, support tickets with names	Internal tools approved for D4 only. External tools need a Data Processing Agreement (DPA) and Board approval.
D5 Employee and special category data	Performance data, sick leave, location tracking, health, religion, union membership	Not allowed in any AI tool without Board approval and Works Council agreement. Always Red.

5. Risk Levels: Green, Amber, Red

Every AI use case is classified into one of three levels. The level decides what happens next. Each rule has an ID so that decisions can cite it, for example "Policy 5.2 A3". **If two rules apply, the highest level wins.**

5.1 Green: allowed now, no review

The employee can start the same day. The decision is logged, but no human review is needed.

Rule	Green criteria
G1	Uses an **internal tool** or an **approved external tool** within its approved data class.
G2	Only D1 Public or D2 Internal data is used.
G3	Output is checked by the employee before use (human in the loop) and is not sent automatically to customers.
G4	No decision is made about a person (customer, user or employee).
G5	The use matches a pre-approved use in KA-BV-AI-2026, Annex 1 (for example drafting, summarising, translating, searching internal knowledge, writing formulas or queries).

Examples: summarising a process guide with the internal assistant; drafting a macro reply template without customer data; generating a spreadsheet formula; translating an internal FAQ.

5.2 Amber: allowed with conditions, Board review

The employee gets the recommendation at once, marked "under review". The case is posted to #ai-governance-review. The Board decides within **5 working days**.

Rule	Amber criteria
A1	Uses a **pilot tool**, or an approved tool outside its listed use.
A2	Uses D3 Confidential data in a tool not explicitly approved for D3.
A3	Uses D4 Personal data (for example customer messages) in an internal tool, at a scale or for a purpose not covered by the tool's listing.
A4	Output goes to customers or users with only sample-based human checks (for example AI-drafted support replies sent after spot checks).
A5	Automates a step that changes how a team works, its workload or its tools in a noticeable way (possible Works Council information duty, KA-BV-AI-2026 section 4).
A6	A public-facing chatbot or AI-generated content that must be labelled as AI under EU AI Act Article 50 transparency rules.

Rule	Amber criteria
A7	The recommendation has low confidence or the sources disagree. When in doubt, Amber.

Examples: AI triage of support tickets that contain user names; testing a pilot meeting-notes tool; an AI assistant on a public help page.

5.3 Red: not allowed until the Board approves

The employee is told clearly that this use is not allowed today and why, and gets the safest allowed alternative if one exists. The case goes to #ai-governance-review for a decision.

Rule	Red criteria
R1	Uses a **not approved tool** or a personal account (for example a private chatbot subscription) for any company data above D1.
R2	Uses D5 Employee or special category data.
R3	**Monitors or evaluates employees**: performance scoring, productivity tracking, behaviour or location analysis. This is a high-risk area under EU AI Act Annex III (employment) and needs Works Council co-determination under BetrVG section 87(1) no. 6.
R4	Supports decisions on hiring, promotion, task allocation based on personal traits, or termination (EU AI Act Annex III, high-risk).
R5	Any practice prohibited by EU AI Act Article 5, for example emotion recognition in the workplace or social scoring.
R6	Fully automated decisions with legal or similar effect on customers or users without human review (GDPR Article 22), for example automatic account blocking without appeal.
R7	D4 Personal data in an external tool without a Data Processing Agreement.

Examples: ranking agents by AI-measured "productivity"; pasting customer chats into a private AI account; using an unlisted browser plugin on support tickets.

6. Internal First: How to Choose a Tool

• **Step 1. Internal first.** Check KA-AI-TOOLS-002 for an internal tool that does the job. Internal tools keep data inside Kleinanzeigen and are supported by the AI Enablement team.

• **Step 2. Approved external second.** If no internal tool fits, use an approved external tool within its data class.

• **Step 3. No match.** If neither exists, do not use an unlisted tool. Submit the need through "Fix My Process" or to #ai-governance-review. Requests with no internal solution are a key input for the AI roadmap.

• **Feedback.** If you choose not to use a recommended internal tool, tell us why. This feedback is how internal tools get better.

7. Human Oversight and Quality

• A named person is responsible for every AI output that is used. AI does not sign, approve or send on its own unless the use is Green under G5 or approved by the Board.

• Check facts, numbers and links in AI output before use. Do not present AI output as verified fact.

• Label AI-generated content that customers or users see, and tell users when they talk to a chatbot (EU AI Act Article 50).

• Keep a source for any recommendation that others will act on.

• Report wrong, biased or harmful outputs to the tool owner.

8. Privacy and GDPR Rules

• **Data minimisation:** enter only the data the task needs. Describe the process, not the people. Do not enter names of colleagues or customers when a description works.

• **Purpose limitation:** do not reuse data in AI tools for a new purpose without checking with the Data Protection Officer.

• **DPIA:** a Data Protection Impact Assessment is required before go-live for every Amber or Red case that involves D4 or D5 data at scale.

• **Retention:** do not store prompts or outputs that contain personal data longer than the task needs.

• **Data subject rights:** tools that process D4 data must support access and deletion requests.

9. Escalation and Review Process

Step	What happens	Time
1. Classify	The use case is classified Green, Amber or Red against section 5 (by "Fix My Process" or by the employee using Annex A).	Same day
2. Escalate	Amber and Red cases are posted to #ai-governance-review with: department (no names), problem, proposed tool, data class, rule IDs that apply.	Same day
3. Review	The Board reviews together at its weekly meeting. Urgent cases: chair plus Data Protection Officer.	Amber: 5 working days; Red: 10 working days
4. Decide	Approve / approve with conditions (for example pilot only, no D4 data) / reject with a reason and an alternative.	With review
5. Works Council	If R3, R4 or A5 applies, the Works Council representative starts the process in KA-BV-AI-2026 section 5.	As per agreement
6. Record	The decision and its reason are logged and the employee can see the final status.	Same day

10. Requesting a New Tool

New tools are requested through "Fix My Process" or #ai-governance-review. A request names the problem, the department, the expected data class, the expected users and why no listed tool fits. Information Security checks the vendor; Legal and the Data Protection Officer check the contract and DPA; the Board decides and, if approved, adds the tool to KA-AI-TOOLS-002 with a tool owner and data class.

11. Incidents and Breaches

If personal or confidential data was entered into a tool that was not approved for it, report it within 24 hours to the Information Security Lead and the Data Protection Officer. Do not try to fix it alone. Reporting in good faith is never a disciplinary matter by itself. Repeated or deliberate breaches of this policy are handled under normal HR processes.

12. Policy Review

The AI Governance Board reviews this policy every 12 months, and earlier when the law, the Approved Tools List or the Works Council Agreement changes. Changes that affect employee monitoring or working conditions are agreed with the Works Council first.

Annex A: Quick Decision Table

Question	If yes	Rule
Does it monitor, score or evaluate employees?	Red	R3
Does it support hiring, promotion, task allocation or termination decisions?	Red	R4
Does it use a not approved tool or a personal account with company data?	Red	R1
Does it use employee or special category data (D5)?	Red	R2
Does it put personal data (D4) into an external tool without a DPA?	Red	R7
Does it decide automatically about customers with legal or similar effect?	Red	R6
Does it use a pilot tool or a tool outside its listed use?	Amber	A1

Question	If yes	Rule
Does it use confidential (D3) or personal (D4) data beyond the tool's listing?	Amber	A2, A3
Does AI output reach customers with only spot checks?	Amber	A4
Does it noticeably change how a team works or its workload?	Amber	A5
Is it a public chatbot or public AI-generated content?	Amber	A6
Is the recommendation uncertain or are sources in conflict?	Amber	A7
None of the above, internal or approved tool, D1 or D2 data, human checks output?	Green	G1 to G5